Repeatability against stopwatch redundancy

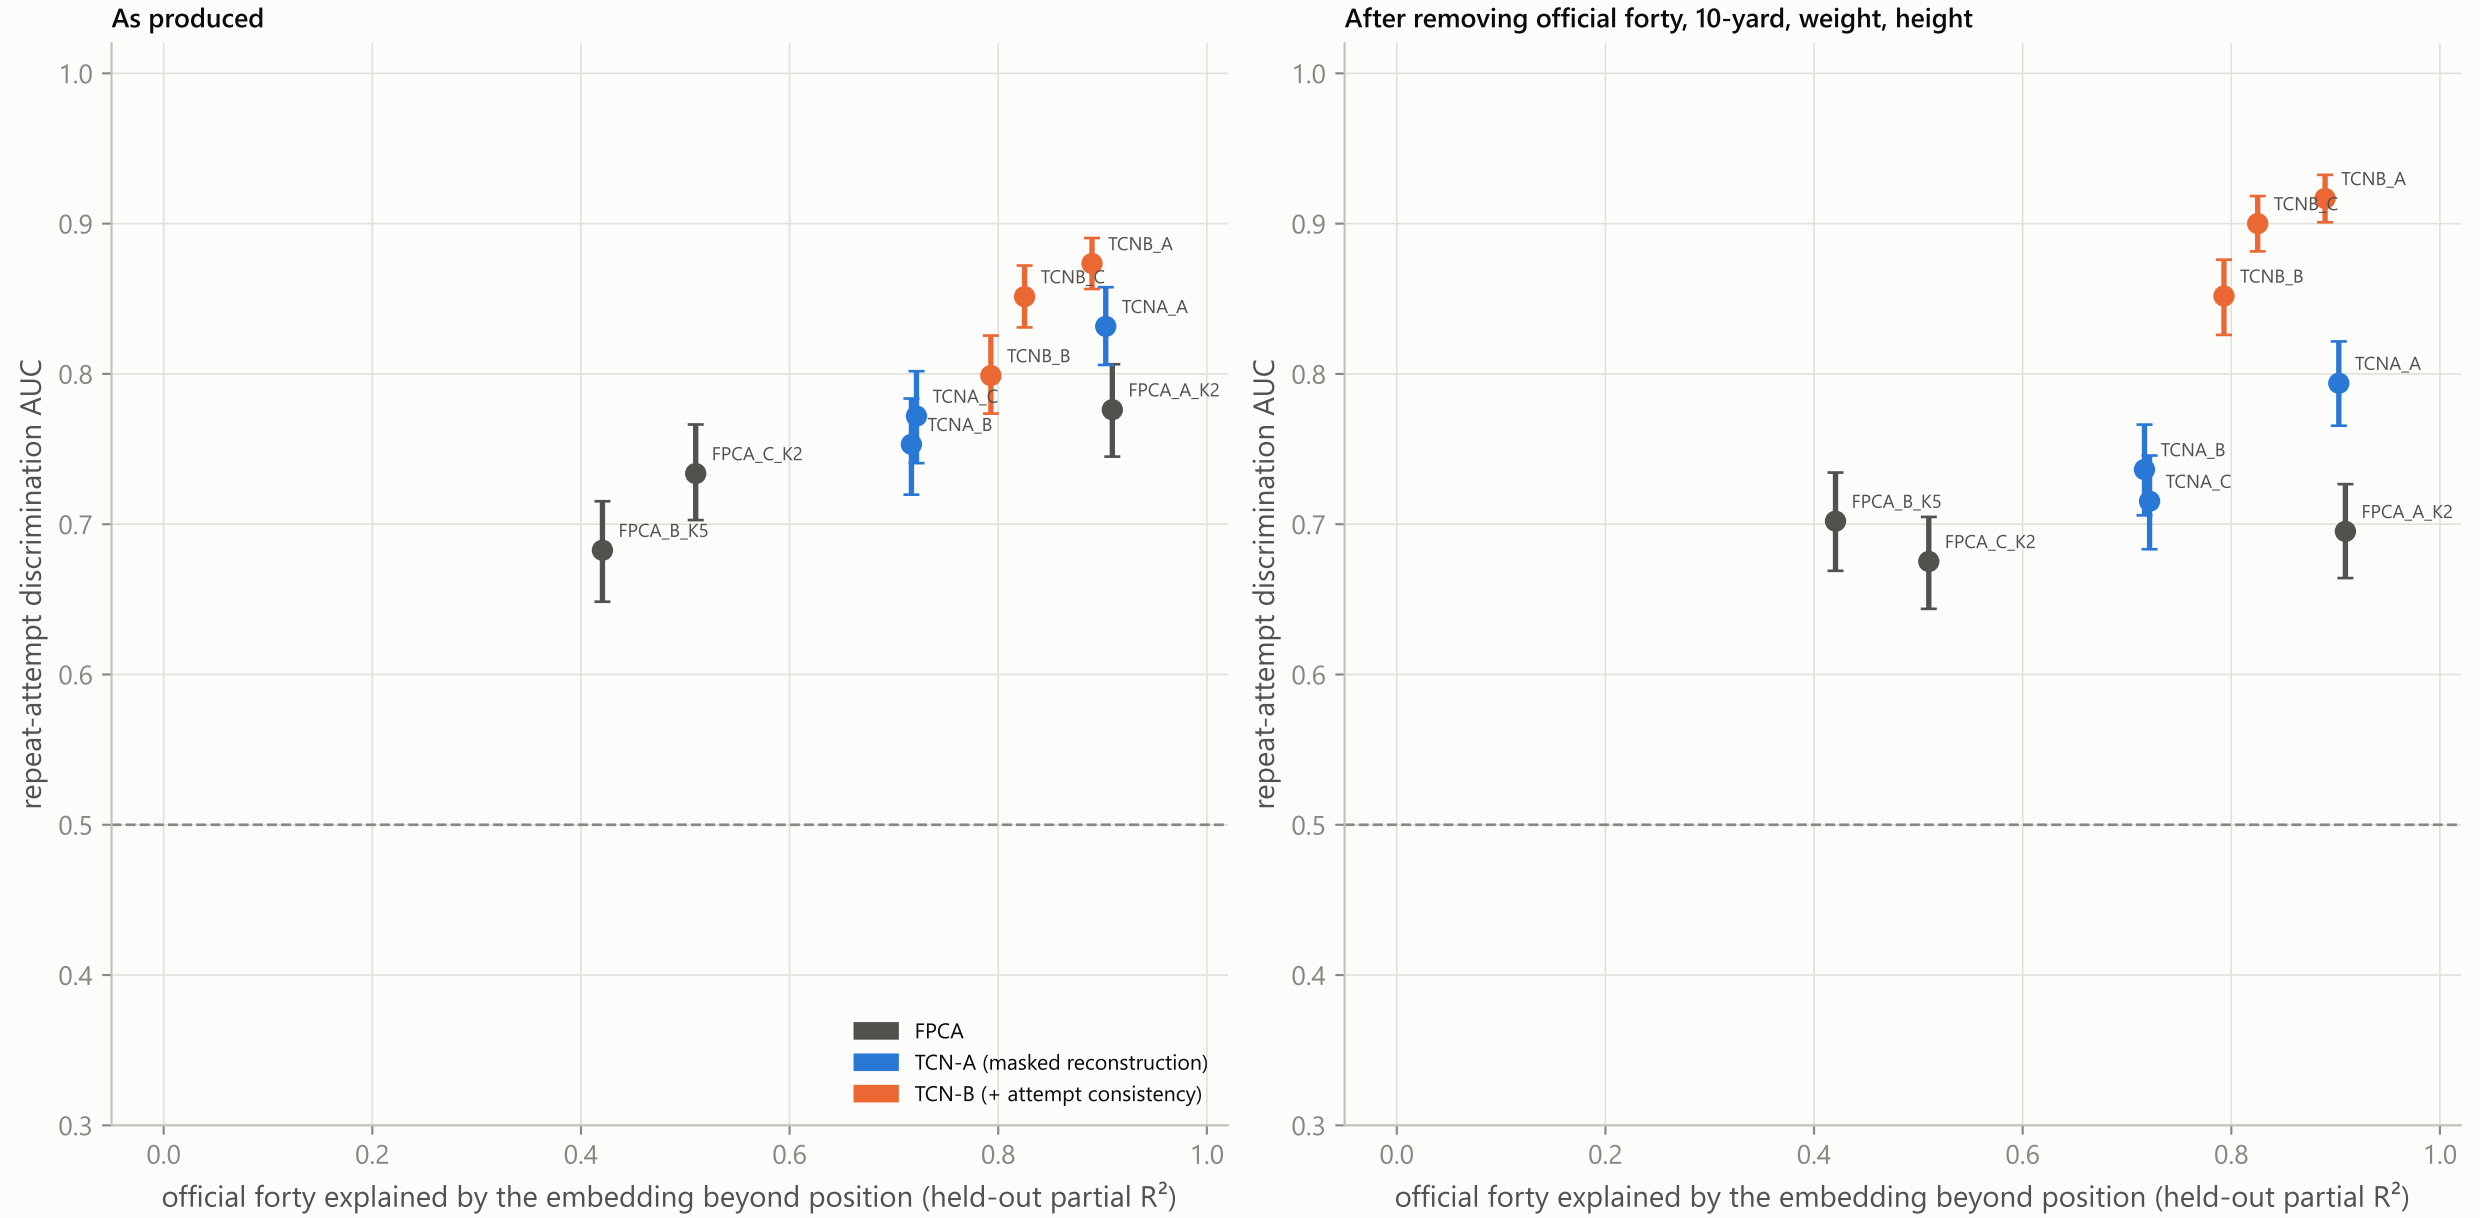

Unit: confirm-role held-out athlete (n up to 265), within position. x is the same on both panels: how much of the official forty the embedding predicts. Left y: can the embedding recognise an athlete from his other attempt? Right y: can it still do so once the cross-fitted linear prediction from the official results is removed? The target region is high on the RIGHT panel; dashed line = chance. Whisker = 95 % athlete bootstrap.