How well does each model predict held-out athletes?

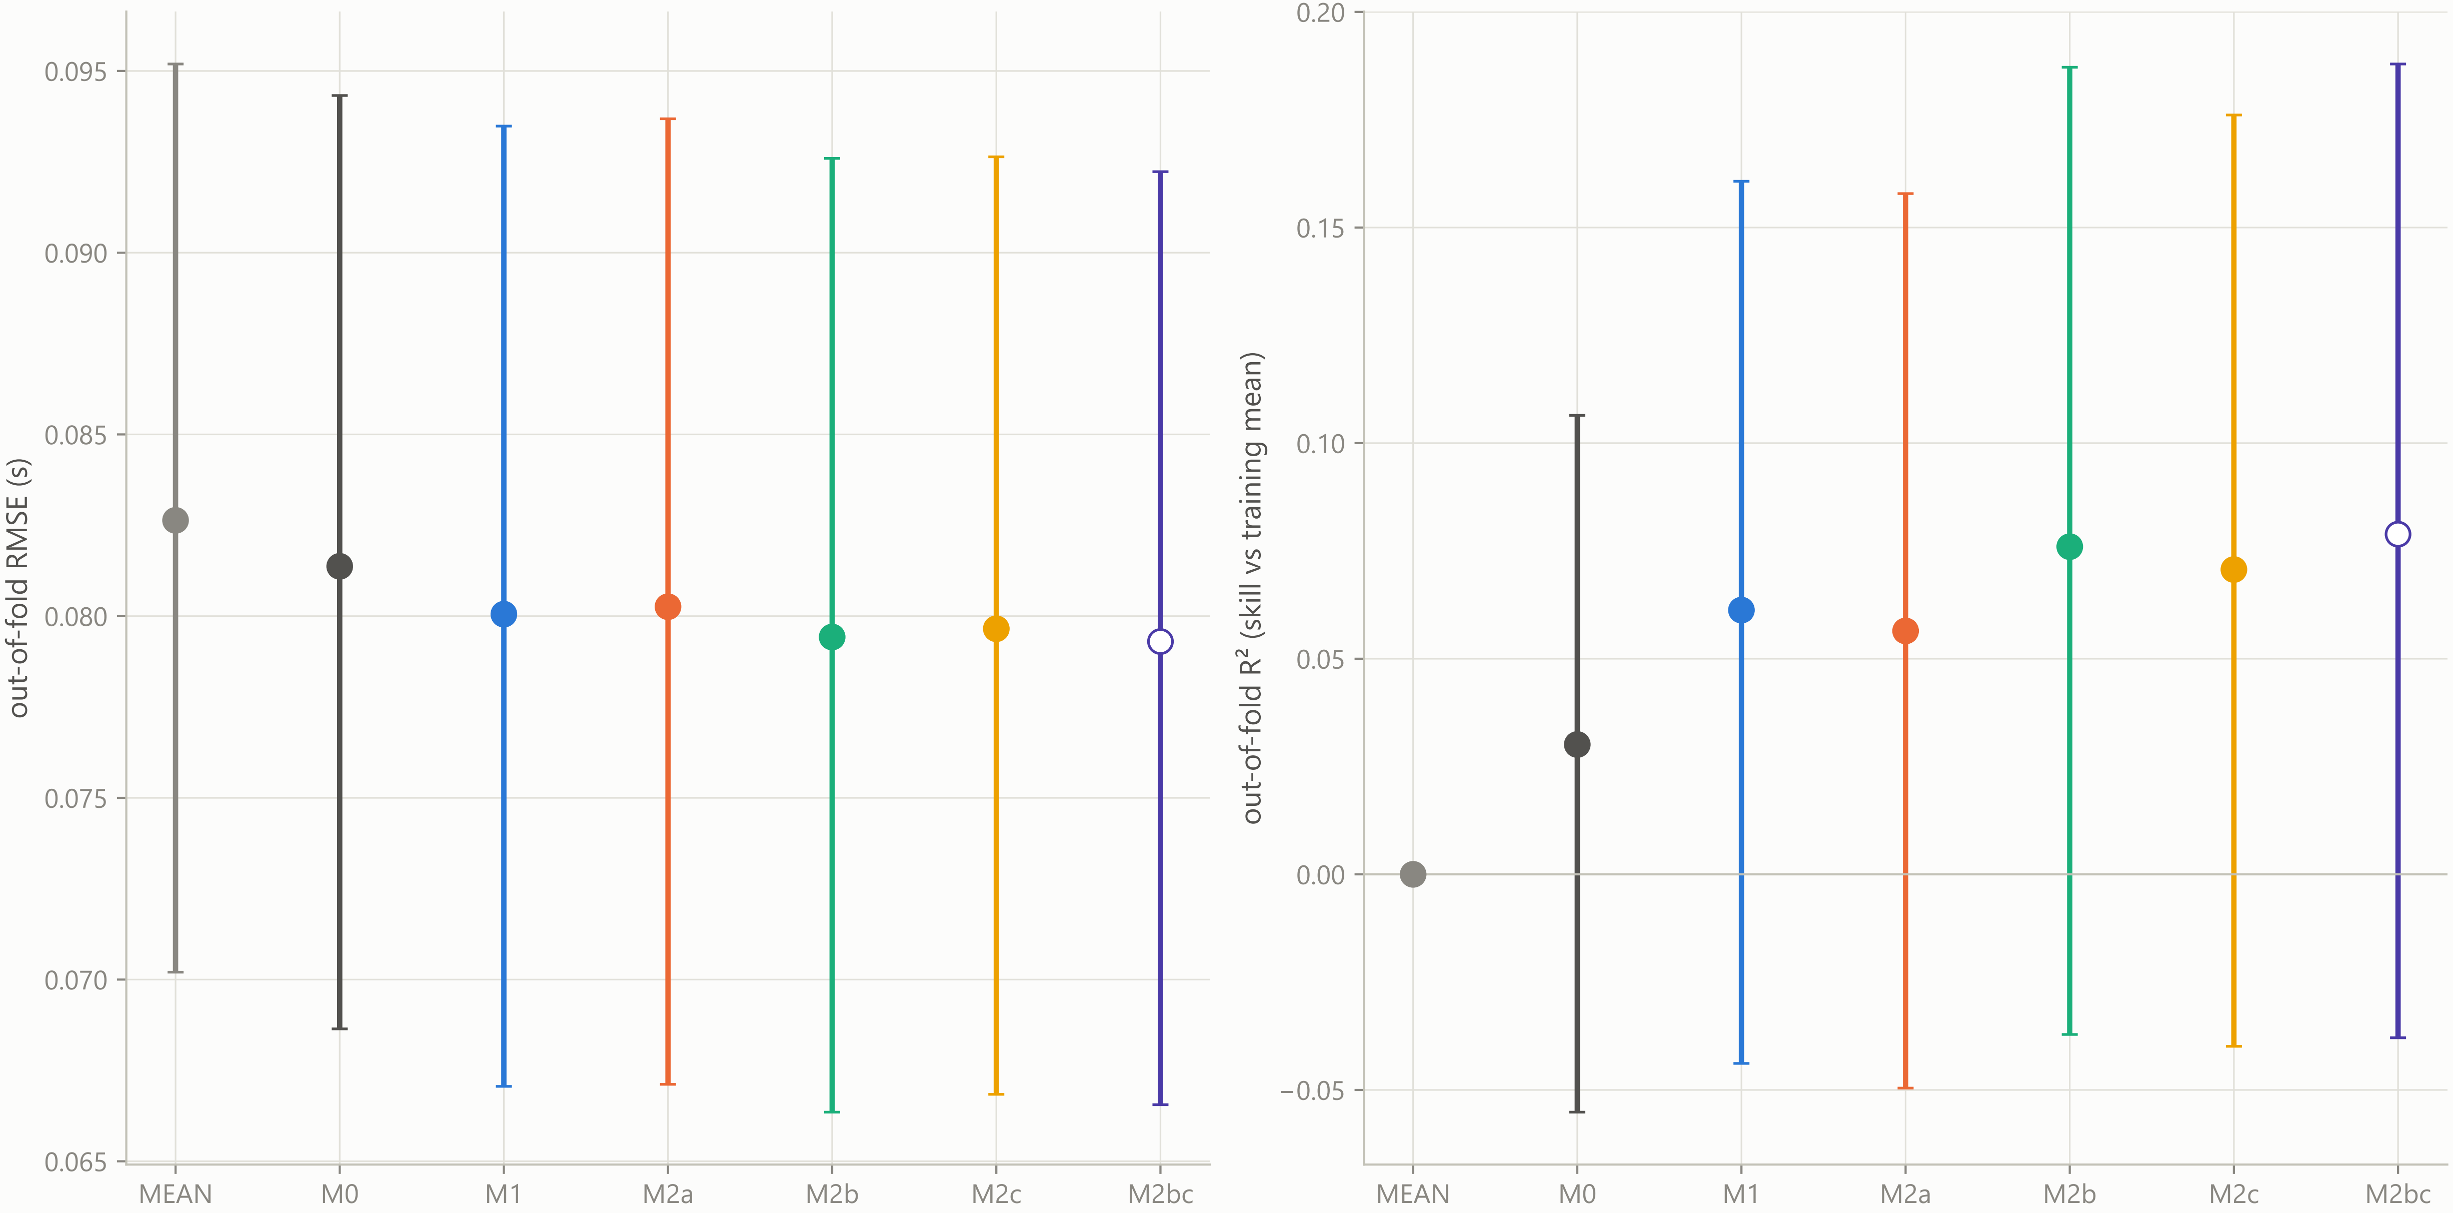

Unit: athlete (n = 64, identical for every model). Repeated athlete-grouped K-fold; preprocessing and the ridge penalty are fitted inside each training fold. Points = mean over repeats, bars = 95 % athlete bootstrap of the out-of-fold predictions. Hollow marker = exploratory model. Overlapping bars do not test a difference: the paired comparison is CN_07. VENDOR DEFINITION UNRESOLVED: the endpoint is an opaque number; no direction or football meaning is implied.